

AI Interview Transcript Analyzer

AI-Driven Interview Understanding, Evaluation
& Decision Intelligence Platform

Project Overview:

The AI Interview Transcript Analyzer is an intelligent end-to-end system designed to transform raw interview audio into structured, actionable HR insights. Built to assist recruiters, HR teams, hiring managers, and talent-acquisition departments, the platform automates the entire interview-analysis workflow from transcription and diarization to sentiment scoring, skill extraction, and AI-powered candidate evaluation.

This project bridges the gap between manual interview review and data-driven decision making by converting unstructured audio into measurable, repeatable, and unbiased evaluation metrics. Powered by Whisper, Resemblyzer, advanced NLP models, Sentence Transformers, and Google Gemini, the system simulates the analysis capabilities of an experienced HR professional, providing structured summaries, skill mappings, sentiment trends, and Hire/No-Hire recommendations.

Scenario 1: HR Recruitment & Candidate Evaluation

Recruiters and hiring managers often spend hours listening to interview recordings, transcribing responses, and making subjective evaluations. The AI Interview Transcript Analyzer streamlines this process by automatically transcribing audio, separating speakers, analysing sentiment, and extracting skills directly from the candidate's speech. For example, an HR team hiring for a customer support role can upload an interview file and instantly receive a polished transcript, emotional tone analysis, and a breakdown of the candidate's strengths.

The system's Gemini-powered evaluation engine generates a concise interview summary, highlights key performance areas, and provides a Hire/No-Hire recommendation with confidence scoring. This ensures assessments remain consistent and objective across all candidates. As a result, organizations can accelerate hiring cycles, reduce manual workload, and base decisions on structured insights rather than subjective impressions.

Scenario 2: Training, Coaching & Interview Preparation

Training centres, mentors, and career coaches often need to review mock interviews and provide detailed feedback to students an effort that becomes time-consuming and repetitive. With the AI Interview Transcript Analyzer, a coach can upload any recorded mock interview and instantly receive a structured transcript, sentiment breakdown, and detected skills, making evaluation faster and more accurate.

The AI-generated performance analysis helps students understand where they excelled and where improvement is needed. Section-wise evaluation, summary insights, and downloadable professional reports allow learners to track progress and prepare more effectively for real interviews. This makes the platform a powerful tool for interview practice, coaching programs, and job readiness preparation.

Architecture Overview:

The AI Interview Transcript Analyzer is built on a modular, multi-layered architecture that connects audio processing, natural language understanding, and intelligent evaluation into a seamless workflow. At its foundation, the system uses Whisper for accurate transcription and Resemblyzer for speaker diarization, enabling the platform to distinguish between interviewers and candidates with high precision. These components work together to transform raw audio into structured, timestamped segments that serve as the basis for further NLP analysis.

On top of the audio layer sits the language intelligence stack, powered by SpaCy, Sentence Transformers, and HuggingFace sentiment models. This layer performs transcript cleaning, sentiment scoring, skills extraction, entity detection, and experience parsing. These insights are then fed into the Gemini LLM, which generates the final structured evaluation including summaries, section-wise performance analysis, and hire/no-hire decisions. Each module functions independently while communicating through centralized pipeline logic, ensuring scalability and clear separation of responsibilities.

The entire system is wrapped in a sleek, interactive Streamlit interface that allows users to upload audio, run analysis, explore dashboards, and download professional reports in TXT, JSON, or PDF. Supporting utilities such as the report exporter, configuration module, and session-state handling ensure smooth UI-backend coordination. This layered architecture enables fast processing, modular extensibility, and a highly maintainable codebase ideal for both enterprise HR systems and training environments.

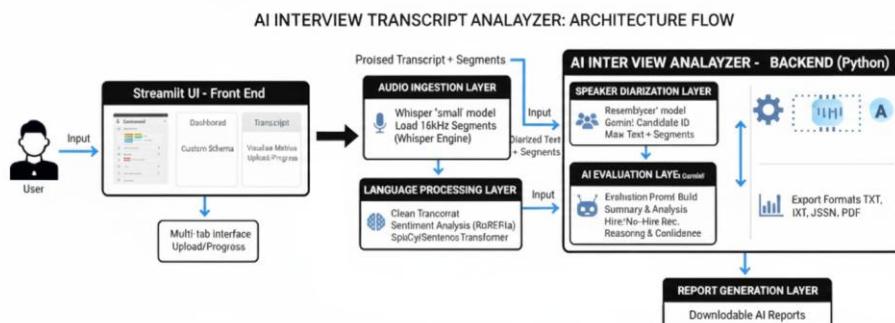


Figure 1: System architecture diagram

Core Technologies:

- **Whisper (Speech-to-Text Engine):**

The system uses OpenAI's Whisper model to perform accurate, multilingual transcription directly from raw audio. It supports segment-level timestamps, enabling precise diarization and downstream language analysis. Whisper ensures high-quality speech recognition even in noisy environments, making the system reliable for real-world interview recordings.

- **Resemblyzer (Speaker Embedding & Diarization):**

Resemblyzer generates voice embeddings that capture speaker identity across audio segments. These embeddings are clustered using Agglomerative Clustering to distinguish between interviewers and candidates, enabling clean separation of speech roles. This supports targeted analysis of candidate responses.

- **SpaCy & Sentence Transformers (Language Processing & Skill Extraction):**

SpaCy performs entity extraction, phrase analysis, and text segmentation, while Sentence Transformers generate semantic embeddings for identifying skills, tools, languages, and academic qualifications. Together, they form the foundation for skill detection and structured candidate profiling.

- **HuggingFace Transformers (Sentiment Analysis):**

A fine-tuned RoBERTa sentiment classifier evaluates emotional tone in the candidate's responses. Segment-level and overall sentiment scores provide insight into confidence, clarity, and communication style throughout the interview.

- **Google Gemini (AI Summary & Hiring Decision):**

Gemini powers the final evaluation module, generating interview summaries, performance breakdowns, and hire/no-hire recommendations.

- **Streamlit (User Interface & Reporting):**

Streamlit forms the front-end layer, delivering an intuitive dashboard for uploading audio, viewing transcript analytics, and downloading reports. Custom CSS styling enhances user experience, while session-state management ensures seamless multi-step analysis.

- **Python Backend & Utilities (Core Logic Layer):**

Python modules manage audio ingestion, metadata extraction, preprocessing, diarization, skill detection, and report generation. Libraries like NumPy, Librosa, SoundFile and ReportLab support efficient audio handling and export functionality. This backend orchestrates all components into a unified processing pipeline.

Component-Wise Architecture:

Component	Description
Streamlit User Interface	Provides an interactive dashboard for uploading interview audio, running AI analysis, viewing dashboards, and downloading reports.
Audio Ingestion Module	Handles audio validation, loading, normalization, and preprocessing using Librosa. Converts audio into consistent formats suitable for Whisper transcription and diarization.
Whisper Transcription Engine	Converts audio into text with timestamped segments. Supports transcription from arrays and file paths, ensuring accurate speech-to-text conversion in various conditions.
Speaker Diarization Module	Uses Resemblyzer embeddings and clustering to identify different speakers. Determines the candidate speaker using Gemini and extracts candidate-specific segments and transcripts.
Transcript Cleaning Module	Removes fillers, normalizes punctuation, and formats segments into readable, structured text. Produces clean transcripts suitable for AI evaluation and report generation.
Sentiment Analysis Engine	Applies a RoBERTa-based sentiment classifier to evaluate emotional tone at both segment and transcript levels, generating detailed sentiment insights.
Skills & Experience Extractor	Uses SpaCy, NLTK, and Sentence Transformer embeddings to detect technical skills, tools, languages, degrees, organizations, and experience durations. Provides structured candidate profiles.
Gemini Evaluation Module	Generates interview summaries, recommendation decisions, performance breakdowns, and reasoning based on all extracted signals transcript, sentiment, skills, and diarized context.
Pipeline Orchestrator	Central controller that coordinates all modules audio ingestion, transcription, diarization, sentiment, skills extraction, and Gemini evaluation into a unified analysis flow.
Report Exporter (TXT/JSON/PDF)	Converts final results into professional reports using JSON encoding and PDF generation with ReportLab. Enables secure export of insights for HR teams and candidates.
Backend Utility & Configuration Layer	Manages API keys, model configurations (Whisper, Gemini, embeddings), skill ontologies, and global constants such as sample rate and file limits.

Pre-requisites:

1. Python Environment Setup

Install Python 3.9+ and create a dedicated virtual environment for clean dependency management.

Official Download: <https://www.python.org/downloads/>

2. Install Required Libraries

All dependencies used in audio processing, NLP, LLM integration, and Streamlit UI must be installed from requirements.txt.

Key libraries include: Whisper, Resemblyzer, Librosa, SpaCy, Sentence Transformers HuggingFace Transformers, Google Generative AI SDK, ReportLab, and Streamlit.

3. Streamlit Installation

Required for running the interactive interview analysis dashboard.

Docs: <https://docs.streamlit.io/library/get-started/installation>

4. NLP & Embedding Models Setup

- SpaCy language model (en_core_web_lg)
- Sentence Transformer model (all-MiniLM-L6-v2)
- HuggingFace sentiment model

SpaCy Models: <https://spacy.io/models>

Sentence Transformers: <https://www.sbert.net/>

5. Google Gemini API Key

Required for summary generation and AI decision-making.

Google Generative AI Setup: <https://ai.google.dev/>

6. Optional Tools for Development

Recommended IDEs:

- Visual Studio Code: <https://code.visualstudio.com/>
- PyCharm Community: <https://www.jetbrains.com/pycharm/>

Project Flow:

1. Environment Setup and Dependency Configuration

- **Activity 1.1:** Obtain and configure the Google Gemini API Key, set up Hugging Face access, and prepare environment variables in .env.
- **Activity 1.2:** Install all required dependencies (Whisper, Resemblyzer, Spacy, HuggingFace models, Streamlit) and validate model loading.
- **Activity 1.3:** Create the project folder structure and configure the config.py file with model names, API keys, and global settings.

2. Core Audio & Language Processing Pipeline Development

- **Activity 2.1:** Implement audio ingestion, format validation, preprocessing, and normalization using Librosa.
- **Activity 2.2:** Integrate the Whisper transcription engine to generate timestamped transcript segments.
- **Activity 2.3:** Implement speaker diarization using Resemblyzer embeddings and clustering; detect candidate speaker and extract candidate-specific segments.
- **Activity 2.4:** Build modules for transcript cleaning, sentiment analysis, skill extraction, and linguistic profiling.

3. AI-Powered HR Evaluation and Decision Logic

- **Activity 3.1:** Engineer evaluation prompts combining transcript, sentiment, skills, and speaker metadata for Gemini.
- **Activity 3.2:** Generate final interview summary, section-wise evaluation, hire/no-hire decision, and reasoning using Gemini.
- **Activity 3.3:** Aggregate all pipeline outputs (transcript, skills, sentiment, metrics) into structured results for UI and reporting.

4. Streamlit UI Implementation and User Interaction

- **Activity 4.1:** Build interactive layout with session state, custom styling, and animated components.

- **Activity 4.2:** Implement multi-tab interface (Dashboard, AI Evaluation, Skills Analysis, Transcript Viewer).
- **Activity 4.3:** Develop the audio upload system, progress indicators, and error handling for the analysis workflow.
- **Activity 4.4:** Display insights and visualizations including transcript segments, skill tags, sentiment metrics, and evaluation results.
- **Activity 4.5:** Implement the report export system to download TXT, JSON, and PDF reports.

5. Testing, Optimization, and Deployment

- **Activity 5.1:** Test UI components, verify audio upload, and validate correct display of metrics, transcript, skills, and evaluation.
- **Activity 5.2:** Conduct deployment preparation by verifying environment variables, optimizing CSS and assets, and ensuring stable end-to-end execution.
- **Activity 5.3:** Perform full pipeline testing on the deployed application, ensuring cross-browser compatibility, responsive UI, and error-free report downloads.

MILESTONE 1: Environment Setup and Dependency Configuration

This foundational milestone establishes the technical environment required for building and deploying the Synthetic Data Factory (SDF). It ensures that all dependencies, frameworks, and integrations are configured correctly for seamless execution of synthetic data generation, visualization, and validation workflows.

Activity 1.1: Obtain Google Gemini API Key

- Sign in or create an account on the official Google AI Studio platform
<https://aistudio.google.com/>

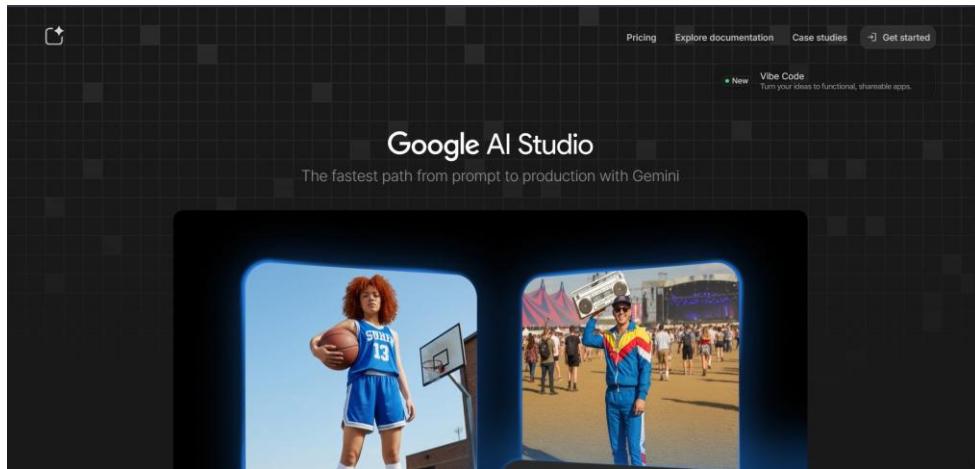


Figure 2: Google AI Studio Console

- Clicking on “Get Started” that’s leads to this Google AI Dashboard:

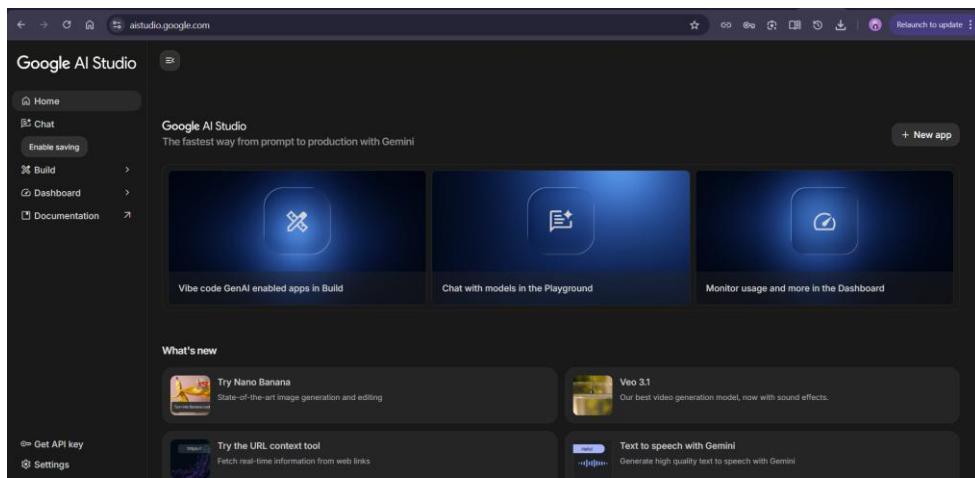


Figure 3: Google AI Studio Dashboard

- Navigate to the Get API Keys section under your account dashboard.

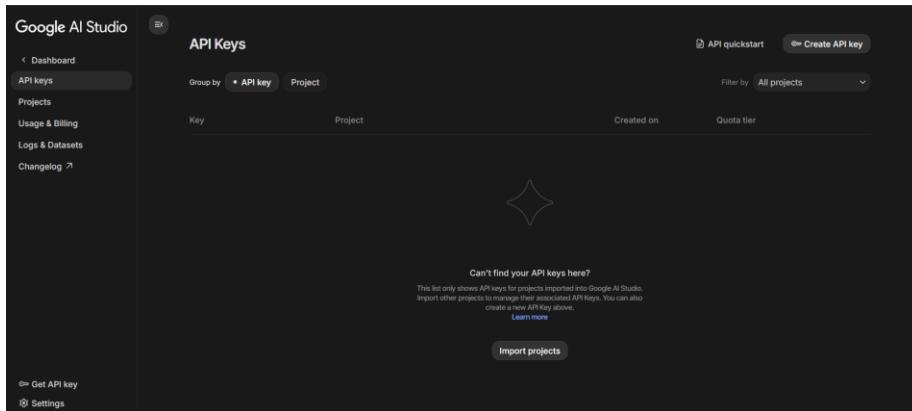


Figure 4: API Key Management Console

- Click on “Create API Key” to generate a new key for your project.
- Name the key appropriately (e.g., AI_Interview_Analyser) for easy identification.

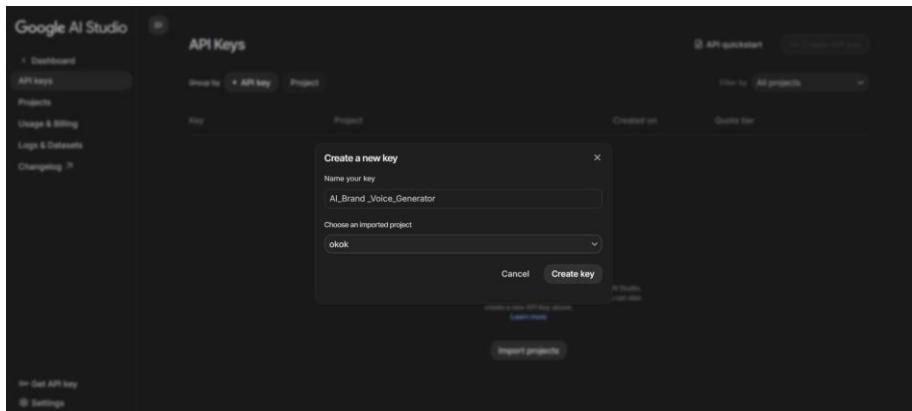


Figure 5: Creating API Key

- Copy and securely store the generated API key for later configuration in your .env file.

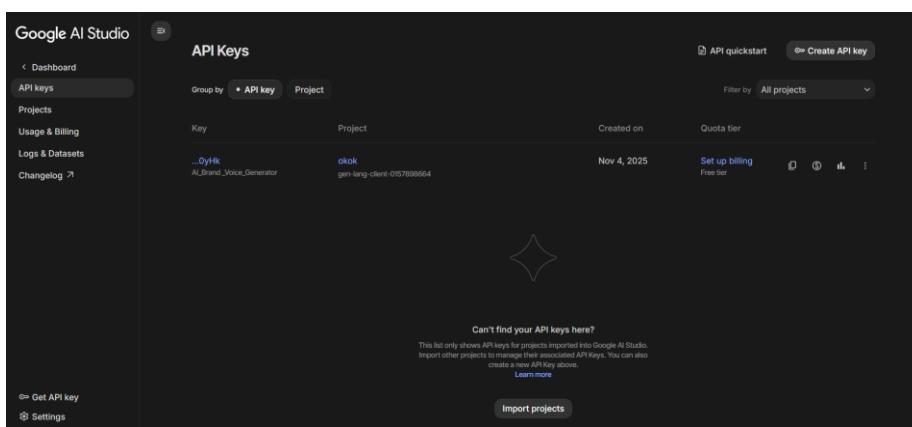


Figure 6: Showing the API Key and Project name on Google Console

Activity 1.2: Set Up Hugging Face Token and Model Access

- Create account or sign in to Hugging Face platform <https://huggingface.co/>

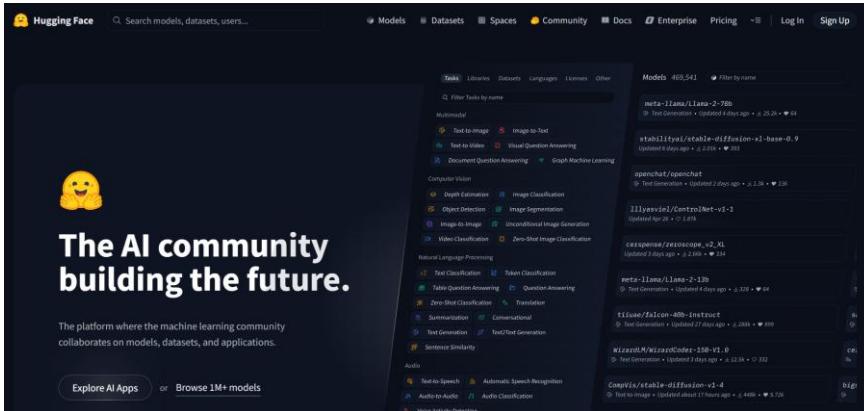


Figure 7: Hugging Face home page

- Navigate to Settings → Access Tokens in your account dashboard

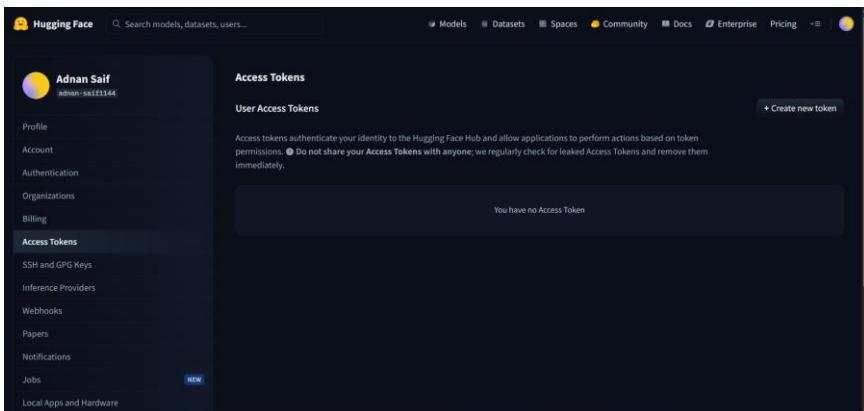


Figure 8: Creating Access Token

- Generate a new token with appropriate permissions for model access

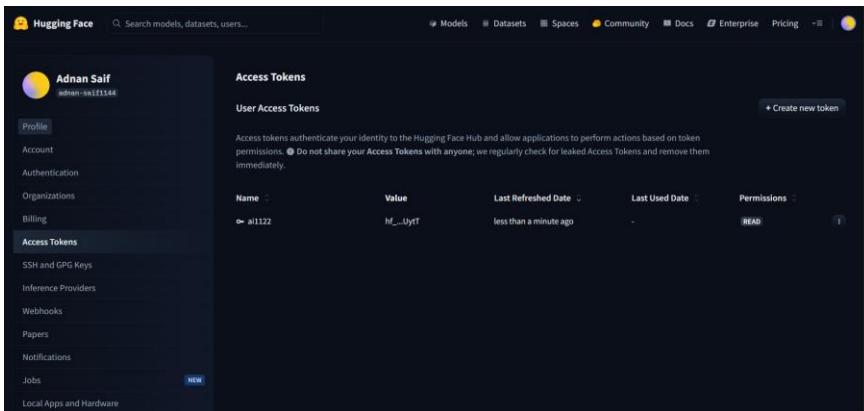


Figure 9: Hugging Face Token

Activity 1.3: Environment Setup & Dependency Installation

- Create a virtual environment for the project:

```
● PS C:\Users\Devansh\Desktop\New folder> python -m venv AIITA
● PS C:\Users\Devansh\Desktop\New folder> AIITA\Scripts\activate
○ (AIITA) PS C:\Users\Devansh\Desktop\New folder> █
```

Figure 10: Creating & Activating Environment

- Install project dependencies:

```
● PS C:\Users\Devansh\Desktop\New folder> python -m venv AIITA
● PS C:\Users\Devansh\Desktop\New folder> AIITA\Scripts\activate
○ (AIITA) PS C:\Users\Devansh\Desktop\New folder> pip install -r requirements.txt█
```

Figure 11: Installing requirements

- Making a .env for securing the Gemini API Key as paste the api key here:

```
2 > ⚙ .env
1   GEMINI_API_KEY = "Your_Gemini_API_Key_Here"
```

Figure 12: .env File

- Configure all environment variables in config.py:
 - GEMINI_API_KEY
 - Model names

```
1 import google.generativeai as genai
2 from dotenv import load_dotenv
3 import os
4
5 load_dotenv()
6 API_KEY = os.getenv("GEMINI_API_KEY")
7
8 genai.configure(api_key = API_KEY)
9
10 # Model Configuration
11 WHISPER_MODEL_SIZE = "small"
12 SENTIMENT_MODEL = "cardiffnlp/twitter-roberta-base-sentiment-latest"
13 EMBEDDING_MODEL = "all-MiniLM-L6-v2"
14 SPACY_MODEL = "en_core_web_lg"
15 GEMINI_MODEL = "gemini-2.0-flash"
16
```

Figure 13: Configure Models

- Set up the project structure for modular development:

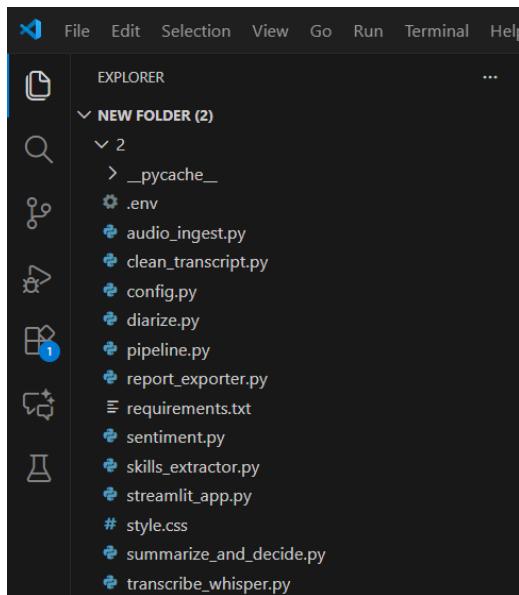


Figure 14: Folder Structure

MILESTONE 2: Core Audio & Language Processing Pipeline Development

This milestone builds the backbone of the analysis system extracting speech, identifying speakers, cleaning transcripts, measuring sentiment, and parsing linguistic details. These components work together to transform raw audio into structured, analysable data.

Activity 2.1: Audio Ingestion & Preprocessing

- Implement audio format validation (.mp3/.wav/.m4a).

```
1 import librosa
2 import numpy as np
3 import soundfile as sf
4 import tempfile
5
6 def validate_audio(file_path):
7     """Validate audio file"""
8     allowed_extensions = ['.mp3', '.wav', '.m4a', '.flac']
9     if not any(file_path.lower().endswith(ext) for ext in allowed_extensions):
10         raise ValueError(f"Unsupported file format. Allowed: {allowed_extensions}")
11 return True
```

Figure 15: Validating audio file

- Load audio with Librosa at 16kHz and normalize amplitude.

```
13 def load_and_preprocess_audio(file_path):
14     """Load and preprocess audio using librosa only"""
15     print("Loading and preprocessing audio...")
16
17     validate_audio(file_path)
18     audio, sr = librosa.load(file_path, sr=16000, mono=True)
19     audio = librosa.util.normalize(audio)
20     print(f"Audio loaded: {len(audio)/sr:.2f} seconds, Sample rate: {sr}")
21
22     return audio, sr
```

Figure 16: Normalize audio

- Save processed audio into temporary WAV files for Whisper.

```
24 def save_audio_to_temp(audio, sr):
25     """Save audio to temporary file for Whisper"""
26     temp_file = tempfile.NamedTemporaryFile(delete=False, suffix='.wav')
27     sf.write(temp_file.name, audio, sr)
28     return temp_file.name
```

Figure 17: Saving processed audio

Activity 2.2: Whisper Transcription Engine Implementation

- Integrate Whisper “small” model for fast transcription.

```

1  import whisper
2  import numpy as np
3  from config import WHISPER_MODEL_SIZE
4
5  # Load Whisper model once
6  whisper_model = whisper.load_model(WHISPER_MODEL_SIZE)

```

Figure 18: Loading Whisper model

- Generate transcript segments with timestamps.

```

8  def transcribe_audio_from_array(audio_array, sample_rate=16000):
9      """Transcribe audio from numpy array - returns segments with timestamps"""
10     print("Transcribing audio with Whisper from array...")
11
12     audio_float = audio_array.astype(np.float32)
13     result = whisper_model.transcribe(audio_float)
14
15     print(f"Transcription complete. Segments: {len(result['segments'])}")
16     return result["segments"], result["text"]

```

Figure 19: Generating Transcript with array

- Validate transcription accuracy on sample audio.

```

18  def transcribe_audio_from_file(audio_path):
19      """Transcribe audio from file path - returns segments with timestamps"""
20      print("Transcribing audio with Whisper from file...")
21
22      result = whisper_model.transcribe(audio_path)
23
24      print(f"Transcription complete. Segments: {len(result['segments'])}")
25      return result["segments"], result["text"]

```

Figure 20: Transcript with audio

Activity 2.3: Speaker Diarization System

- Extract speaker embeddings using Resemblyzer.

```

1  import numpy as np
2  import librosa
3  from resemblyzer import VoiceEncoder, preprocess_wav
4  from sklearn.cluster import AgglomerativeClustering
5  from sklearn.metrics import silhouette_score
6  from config import WHISPER_MODEL_SIZE
7  import google.generativeai as genai
8  from config import GEMINI_MODEL
9
10 model = genai.GenerativeModel(GEMINI_MODEL)
11
12 # Load models
13 encoder = VoiceEncoder()
14
15 def get_segment_embedding_from_array(audio_array, sample_rate, start, end):
16     """Get embedding from audio array segment"""
17     segment = audio_array[int(start * sample_rate):int(end * sample_rate)]
18     segment = preprocess_wav(segment)
19     embed = encoder.embed_utterance(segment)
20     return embed

```

Figure 21: Embedding with resemblyzer

- Perform clustering using Agglomerative Clustering.

```

22 def detect_num_speakers(embeddings, max_speakers=4):
23     """Your existing speaker detection"""
24     best_score = -1
25     best_k = 1
26
27     for k in range(2, max_speakers + 1):
28         clustering = AgglomerativeClustering(n_clusters=k).fit(embeddings)
29         score = silhouette_score(embeddings, clustering.labels_)
30         if score > best_score:
31             best_score = score
32             best_k = k
33     return best_k
34
35 def diarize_whisper_segments_from_array(audio_array, sample_rate, whisper_segments, max_speakers=4):
36     """Diarization using audio array instead of file path"""
37     embeddings = []
38
39     print("Generating speaker embeddings from audio array...")
40     for seg in whisper_segments:
41         emb = get_segment_embedding_from_array(audio_array, sample_rate, seg["start"], seg["end"])
42         embeddings.append(emb)
43
44     if not embeddings:
45         return []
46
47     embeddings = np.vstack(embeddings)
48
49     num_speakers = detect_num_speakers(embeddings, max_speakers)
50     print(f"Detected speakers: {num_speakers}")
51
52     clustering = AgglomerativeClustering(n_clusters=num_speakers)
53     labels = clustering.fit_predict(embeddings)
54
55     diarized = []
56     for seg, spk in zip(whisper_segments, labels):
57         diarized.append({
58             "speaker": f"Speaker_{spk+1}",
59             "start": seg["start"],
60             "end": seg["end"],
61             "text": seg["text"]
62         })
63
64     return diarized
65

```

Figure 22: Clustering the speaker

- Auto-detect candidate speaker using Gemini by analysing early conversation segments.

```

66     def determine_candidate_speaker(diarized_segments):
67         """Use Gemini to determine which speaker is the candidate/interviewee"""
68         try:
69             conversation_text = ""
70             for segment in diarized_segments[:10]:
71                 conversation_text += f"{segment['speaker']}: {segment['text']}\n"
72
73             prompt = f"""
74             Analyze this interview conversation and identify which speaker is the candidate/interviewee (the person being interviewed).
75
76             Conversation:
77             {conversation_text}
78
79             Based on the content, which speaker is most likely the candidate being interviewed?
80             Return ONLY the speaker name exactly as it appears in the conversation (e.g., "speaker_1" or "Speaker_2").
81             Do not add any explanation, only the speaker name as mention in conversation.
82             """
83
84             response = model.generate_content(prompt)
85             candidate_speaker = response.text.strip()
86
87             return candidate_speaker
88
89         except Exception:
90             return "Speaker_1"

```

Figure 23: Detection of Candidate

- Candidate diarized segments with Whisper timestamps.

```

92     def get_candidate_segments(diarized_segments, candidate_speaker="Speaker_1"):
93         """Extract only candidate segments"""
94         return [seg for seg in diarized_segments if seg.get("speaker") == candidate_speaker]
95
96     def get_candidate_transcript(diarized_segments, candidate_speaker="Speaker_1"):
97         """Get candidate's full transcript"""
98         candidate_segments = get_candidate_segments(diarized_segments, candidate_speaker)
99         return " ".join([seg.get("text", "") for seg in candidate_segments])

```

Figure 24: Extracting Candidate Segments

Activity 2.4: Transcript Cleaning, Sentiment Processing & Skill Extraction

- Remove fillers, normalize punctuation, and structure dialogue.

```

1 import re
2 import nltk
3 from nltk.tokenize import sent_tokenize
4
5 nltk.download('punkt', quiet=True)
6
7 def clean_text(text):
8     """Clean transcript text - remove fillers, normalize punctuation"""
9     fillers = [
10         "uh", "um", "uhh", "umm",
11         "ah", "er", "eh", "hmm",
12         "kinda", "sorta",
13         "oh", "ohhh",
14         "huh", "mm-hmm", "mmhm",
15         "lmao", "lol"
16     ]
17
18     for filler in fillers:
19         text = re.sub(r'(b' + filler + r'\b)', '', text, flags=re.IGNORECASE)
20
21     text = re.sub(r'\s+', ' ', text)
22
23     text = re.sub(r'([.,!?:;])', r'\1 ', text)
24     text = re.sub(r'([.,!?:;])(\w)', r'\1 \2', text)
25
26     return text.strip()
27
28 def clean_transcript_segments(segments):
29     """Clean all transcript segments"""
30     cleaned_segments = []
31     for segment in segments:
32         cleaned_text = clean_text(segment.get('text', ''))
33         if cleaned_text:
34             cleaned_segments.append({
35                 '+segment',
36                 'text': cleaned_text
37             })
38     return cleaned_segments
39
40 def format_transcript_for_display(segments):
41     """Format transcript for nice display"""
42     formatted = []
43     for seg in segments:
44         formatted.append(f'{seg["speaker"]}: {seg["start"]:.2f}s → {seg["end"]:.2f}s\n{seg["text"]}\n')
45     return '\n'.join(formatted)

```

Figure 25: Filtering the Text

- Perform segment-level and full-transcript sentiment analysis.

```

1 from transformers import AutoTokenizer, AutoModelForSequenceClassification
2 import torch
3 from config import SENTIMENT_MODEL
4
5 # Load model once
6 sent_model = AutoModelForSequenceClassification.from_pretrained(SENTIMENT_MODEL)
7 tokenizer = AutoTokenizer.from_pretrained(SENTIMENT_MODEL)
8
9 def analyze_sentiment(text):
10     """Your existing sentiment analysis function"""
11     text = text.replace("\n", " ")
12
13     inputs = tokenizer(
14         text,
15         return_tensors="pt",
16         truncation=True,
17         max_length=512,
18         padding="max_length",
19     )
20
21     inputs.pop("token_type_ids", None)
22
23     with torch.no_grad():
24         outputs = sent_model(**inputs)
25
26     scores = torch.softmax(outputs.logits, dim=1)[0].tolist()
27     labels = ["negative", "neutral", "positive"]
28
29     max_score = max(scores)
30     max_index = scores.index(max_score)
31
32     return {
33         "label": labels[max_index],
34         "score": max_score,
35         "scores": {
36             "negative": scores[0],
37             "neutral": scores[1],
38             "positive": scores[2]
39         }
40     }
41
42 def analyze_segment_sentiments(segments):
43     """Analyze sentiment for each segment"""
44     segment_sentiments = []
45     for segment in segments:
46         sentiment = analyze_sentiment(segment['text'])
47         segment_sentiments.append({
48             '+segment',
49             "sentiment": sentiment
50         })
51     return segment_sentiments

```

Figure 26: Segmentation

- Use SpaCy + SentenceTransformer to extract:
 - Skills
 - Tools
 - Languages
 - Education
 - Organizations

```

51     # Phrase-level embedding matching
52     sentences = nltk.sent_tokenize(text)
53     all_phrases = set()
54
55     for sent in sentences:
56         words = sent.split()
57         for n in [1, 2, 3]:
58             for gram in ngrams(words, n):
59                 phrase = " ".join(gram)
60                 all_phrases.add(phrase)
61
62     phrase_list = list(all_phrases)
63     if phrase_list:
64         phrase_embeddings = emb_model.encode(phrase_list, convert_to_tensor=True)
65         sim = util.cos_sim(phrase_embeddings, skill_embeddings)
66
67         for i, phrase in enumerate(phrase_list):
68             score = float(sim[i].max())
69             idx = int(sim[i].argmax())
70
71             if score > 0.60:
72                 extracted["skills"].add(MASTER_SKILLS[idx])
73
74     # Language extraction
75     for lang in LANGUAGE_SKILLS:
76         if lang.lower() in text.lower():
77             extracted["languages"].add(lang)
78
79     # Degree extraction
80     for deg in DEGREES:
81         if deg.lower() in text.lower():
82             extracted["degrees"].add(deg)
83
84     # Convert sets to lists for JSON serialization
85     for key in extracted:
86         if isinstance(extracted[key], set):
87             extracted[key] = list(extracted[key])
88
89     return extracted
90
91
92 def format_experience_durations(durations):
93     """Format experience durations for display"""
94     formatted = []
95     for duration in durations:
96         if isinstance(duration, dict):
97             formatted.append(f"{duration.get('value', '')} {duration.get('unit', '')}")
98         else:
99             formatted.append(str(duration))
100    return formatted

```

Figure 27: Skill Extraction

```

51     # Phrase-level embedding matching
52     sentences = nltk.sent_tokenize(text)
53     all_phrases = set()
54
55     for sent in sentences:
56         words = sent.split()
57         for n in [1, 2, 3]:
58             for gram in ngrams(words, n):
59                 phrase = " ".join(gram)
60                 all_phrases.add(phrase)
61
62     phrase_list = list(all_phrases)
63     if phrase_list:
64         phrase_embeddings = emb_model.encode(phrase_list, convert_to_tensor=True)
65         sim = util.cos_sim(phrase_embeddings, skill_embeddings)
66
67         for i, phrase in enumerate(phrase_list):
68             score = float(sim[i].max())
69             idx = int(sim[i].argmax())
70
71             if score > 0.60:
72                 extracted["skills"].add(MASTER_SKILLS[idx])
73
74     # Language extraction
75     for lang in LANGUAGE_SKILLS:
76         if lang.lower() in text.lower():
77             extracted["languages"].add(lang)
78
79     # Degree extraction
80     for deg in DEGREES:
81         if deg.lower() in text.lower():
82             extracted["degrees"].add(deg)
83
84     # Convert sets to lists for JSON serialization
85     for key in extracted:
86         if isinstance(extracted[key], set):
87             extracted[key] = list(extracted[key])
88
89     return extracted
90
91
92 def format_experience_durations(durations):
93     """Format experience durations for display"""
94     formatted = []
95     for duration in durations:
96         if isinstance(duration, dict):
97             formatted.append(f"{duration.get('value', '')} {duration.get('unit', '')}")
98         else:
99             formatted.append(str(duration))
100    return formatted

```

Figure 28: Skill Extraction

MILESTONE 3: AI-Powered HR Evaluation & Decision Engine

This milestone implements the high-level reasoning capabilities that convert raw analysis into meaningful HR insights. Gemini generates summaries, performance breakdowns, and hiring decisions based on structured data.

Activity 3.1: AI Evaluation Prompt Engineering

- Build a structured, context-rich evaluation prompt including:
 - Transcript
 - Sentiment result
 - Skills

```

2 > ⚡ summarize_and_decide.py > ⏪ generate_evaluation
1  import google.generativeai as genai
2  from config import GEMINI_MODEL
3
4  model = genai.GenerativeModel(GEMINI_MODEL)
5
6  def build_evaluation_prompt(transcript, sentiment, skills_info, diarized_segments):
7      """Build comprehensive prompt for Gemini evaluation"""
8      prompt = f"""
9      You are an HR interview evaluation assistant.
10
11     INTERVIEW TRANSCRIPT:
12     {transcript}
13
14     SENTIMENT ANALYSIS: {sentiment['label']} (confidence: {sentiment['score']:.2f})
15
16     SKILLS DETECTED:
17     - Technical Skills: {', '.join.skills_info['skills']} if skills_info['skills'] else 'None'
18     - Languages: {', '.join.skills_info['languages']} if skills_info['languages'] else 'None'
19     - Tools: {', '.join.skills_info['tools']} if skills_info['tools'] else 'None'
20     - Education: {', '.join.skills_info['degrees']} if skills_info['degrees'] else 'None'
21     - Organizations: {', '.join.skills_info['organizations']} if skills_info['organizations'] else 'None'
22     - Experience: {skills_info['experience_durations']} if skills_info['experience_durations'] else 'None'
23
24     YOUR TASKS:
25     1. Generate a concise summary of the interview (4-5 lines)
26     2. Provide a clear HIRE or NO-HIRE recommendation
27     3. Provide confidence score (0-100)
28     4. Divide interview into 3 logical sections and evaluate candidate performance in each
29     5. Mention key points discussed in each section
30     6. Provide reasoning for your recommendation

```

Figure 29: AI Evaluation

- Define output format for Gemini:

```

32     OUTPUT FORMAT:
33     Summary: [4-5 line summary]
34
35     Recommendation: [HIRE/NO-HIRE]
36
37     Confidence: [0-100]
38
39     Section 1: [Section name]
40     - Performance: [Evaluation]
41     - Key Points: [Key discussion points]
42
43     Section 2: [Section name]
44     - Performance: [Evaluation]
45     - Key Points: [Key discussion points]
46
47     Section 3: [Section name]
48     - Performance: [Evaluation]
49     - Key Points: [Key discussion points]
50
51     Reasoning: [Detailed reasoning for recommendation]
52
53     Do not use markdown or bold formatting, provide proper spacing for paragraph and lines.
54     """
55     return prompt

```

Figure 30: Gemini Output Format

Activity 3.2: Generate Final Evaluation with Gemini

- Generate professional summary of the interview.

```

57 def generate_evaluation(transcript, sentiment, skills_info, diarized_segments):
58     """Generate final evaluation using Gemini"""
59     prompt = build_evaluation_prompt(transcript, sentiment, skills_info, diarized_segments)
60
61     try:
62         response = model.generate_content(prompt)
63         return response.text
64     except Exception as e:
65         return f"Error generating evaluation: {str(e)}"

```

Figure 31: Generating Output with Gemini

- Produce section-wise evaluation across three logical interview phases.
- Generate Hire/No-Hire recommendation with reasoning.
- Store results for reporting and UI display.

Activity 3.3: Final Pipeline

- Combine sentiment, skills, transcript, and speaker metrics.

```

1  from audio_ingest import load_and_preprocess_audio
2  from transcribe_whisper import transcribe_audio_from_array
3  from diarize import diarize_whisper_segments_from_array, get_candidate_transcript, get_candidate_segments, determine_candidate
4  from clean_transcript import clean_transcript_segments, format_transcript_for_display
5  from sentiment import analyze_sentiment, analyze_segment_sentiments
6  from skills_extractor import extract_candidate_info
7  from summarize_and_decide import generate_evaluation
8
9  def run_full_pipeline(audio_path, candidate_speaker=None):
10     """Run the complete interview analysis pipeline"""
11
12     print("Starting Interview Analysis Pipeline...")
13
14     # Audio ingestion
15     print("Loading audio...")
16     audio_array, sample_rate = load_and_preprocess_audio(audio_path)
17
18     # Transcription
19     print("Transcribing audio...")
20     whisper_segments, full_transcript = transcribe_audio_from_array(audio_array, sample_rate)
21
22     # Diarization
23     print("Speaker diarization...")
24     diarized_segments = diarize_whisper_segments_from_array(audio_array, sample_rate, whisper_segments)
25
26     # Determine candidate speaker
27     if candidate_speaker is None:
28         print("Determining candidate speaker using AI...")
29         candidate_speaker = determine_candidate_speaker(diarized_segments)
30         print(f"Detected candidate speaker: {candidate_speaker}")
31
32     # Transcript cleaning
33     print("Cleaning transcript...")
34     cleaned_segments = clean_transcript_segments(diarized_segments)

```

Figure 32: Pipeline

```

32     # Transcript cleaning
33     print("Cleaning transcript...")
34     cleaned_segments = clean_transcript_segments(diarized_segments)
35
36     # Get candidate transcript
37     print("Extracting candidate speech...")
38     candidate_segments = get_candidate_segments(cleaned_segments, candidate_speaker)
39     candidate_transcript = get_candidate_transcript(cleaned_segments, candidate_speaker)
40
41     # Sentiment analysis
42     print("Analyzing sentiment...")
43     sentiment = analyze_sentiment(candidate_transcript)
44     segment_sentiments = analyze_segment_sentiments(cleaned_segments)
45
46     # Skills extraction
47     print("Extracting skills...")
48     skills_info = extract_candidate_info(candidate_transcript)
49
50     # AI Evaluation
51     print("Generating AI evaluation...")
52     formatted_transcript = format_transcript_for_display(cleaned_segments)
53     evaluation = generate_evaluation(
54         formatted_transcript,
55         sentiment,
56         skills_info,
57         cleaned_segments
58     )
59

```

Figure 33: Pipeline

- Compute final metadata including:
 - Duration
 - Segment count
 - Speaker distribution

```

60     # Compile final results
61     results = {
62         'audio_metadata': {
63             'duration': len(audio_array)/sample_rate,
64             'sample_rate': sample_rate
65         },
66         'full_transcript': full_transcript,
67         'diarized_segments': cleaned_segments,
68         'candidate_segments': candidate_segments,
69         'candidate_transcript': candidate_transcript,
70         'sentiment': sentiment,
71         'segment_sentiments': segment_sentiments,
72         'skills_info': skills_info,
73         'evaluation': evaluation,
74         'candidate_speaker': candidate_speaker
75     }
76
77     print("Pipeline completed successfully!")
78     return results

```

Figure 34: Storing result

MILESTONE 4: Streamlit UI Integration & Application Logic

This milestone focuses on developing a fully interactive, visually refined, and user-friendly platform that brings together every backend component of the AI Interview Transcript Analyzer into a seamless, cohesive user experience. It ensures that the complex processes of audio ingestion, transcription, diarization, sentiment analysis, skills extraction, and AI-driven evaluation are presented to the user in an intuitive, well-structured interface. Through a polished Streamlit UI enhanced with custom styling, responsive design, and multi-tab navigation, this milestone transforms the underlying AI pipeline into a professional, easy-to-use application that allows users to upload interview audio, monitor analysis progress in real time, explore insights, and download reports all within a smooth and engaging workflow.

Activity 4.1: Interactive Layout & Session Management

- Implement session state for:
 - Uploaded file
 - Analysis results
 - UI navigation

```

2 > ♦ streamlit_app.py > ...
  1  import streamlit as st
  2  import tempfile
  3  import os
  4  from datetime import datetime
  5  from pipeline import run_full_pipeline
  6  from report_exporter import export_txt, export_json, export_pdf
  7  import json
  8  import base64
  9
 10 # Page configuration
 11 st.set_page_config(
 12     page_title="AI Interview Transcript Analyzer",
 13     layout="wide",
 14     initial_sidebar_state="collapsed"
 15 )
 16
 17 # Load CSS
 18 def load_css():
 19     with open("2/style.css") as f:
 20         st.markdown(f"<style>{f.read()}</style>", unsafe_allow_html=True)
 21
 22 def get_base64_of_bin_file(bin_file):
 23     with open(bin_file, "rb") as f:
 24         data = f.read()
 25     return base64.b64encode(data).decode()
 26
 27 def set_background(image_file):
 28     bin_str = get_base64_of_bin_file(image_file)
 29     css = f"""
 30     <style>
 31         .stApp {{
 32             background-image: url("data:image/png;base64,{bin_str}");
 33             background-size: cover;
 34             background-attachment: fixed;
 35         }}
 36     </style>
 37     """
 38     st.markdown(css, unsafe_allow_html=True)
 39
 40 def main():
 41     # Load custom CSS
 42     load_css()
 43
 44     st.markdown("""
 45         <div class="main-header fade-in">
 46             <h1 style="margin:0; font-size: 3rem; font-weight: 700;">AI Interview Transcript Analyzer</h1>
 47             <p style="margin:0; font-size: 1.2rem; opacity: 0.9;">Professional Interview Analysis Platform</p>
 48         </div>
 49     """, unsafe_allow_html=True)
 50
 51     # Initialize session state
 52     if 'analysis_results' not in st.session_state:
 53         st.session_state.analysis_results = None
 54     if 'analysis_complete' not in st.session_state:
 55         st.session_state.analysis_complete = False
 56     if 'uploaded_file' not in st.session_state:
 57         st.session_state.uploaded_file = None
 58

```

Figure 35: State Management & Streamlit Configuration

- Configure animations, transitions, and custom CSS.

```
> /* style.css > */
1 /* Global Styles */
2 .main {
3     background: linear-gradient(135deg, #f667eaa 0%, #f764ba2 100%);
4 }
5
6 #app {
7     background: linear-gradient(135deg, #f5f7fa 0%, #c3cfec 100%);
8 }
9
10 /* Header Styles */
11 .main-header {
12     text-align: center;
13     padding: 2rem 0;
14     background: linear-gradient(135deg, #f667eaa 0%, #f764ba2 100%);
15     color: #fff;
16     border-radius: 15px;
17     border-radius: 15px;
18     margin-bottom: 2rem;
19     box-shadow: 0 10px 30px rgba(0,0,0,0.2);
20     position: relative;
21     z-index: 1;
22     overflow: hidden;
23 }
24
25 .main-header::before {
26     content: '';
27     position: absolute;
28     top: 0;
29     left: 0;
30     right: 0;
31     bottom: 0;
32     background: url('data:image/svg+xml,<svg xmlns="http://www.w3.org/2000/svg" viewBox="0 0 100 100"><defs><pattern id="grid" width="10" height="10" patternUnits="userSpaceOnUse"><path d="M 10 0 L 0 0 0 10" fill="none" stroke="rgb(255,255,255,0.2)">/</path></pattern></defs><g><rect fill="white" width="100" height="100"/><rect fill="#f667eaa" width="100" height="100" pattern="url(#grid)"/></g></svg>') no-repeat;
33 }
34
35 /* Card Styles */
36 .custom-card {
37     background: #fff;
38     padding: 1.5rem;
39     border-radius: 15px;
40     box-shadow: 0 5px 20px rgba(0,0,0,0.2);
41     border: 1px solid #f667eaa;
42     transition: all 0.3s ease;
43     margin-bottom: 1rem;
44 }
45
46 .custom-card:hover {
47     transform: translateY(-5px);
48     box-shadow: 0 15px 35px rgba(0,0,0,0.15);
49 }
50
51 /* Metric Cards */
52 .metric-card {
53     background: linear-gradient(135deg, #f667eaa 0%, #f764ba2 100%);
54     color: #fff;
55     padding: 1.5rem;
56     border-radius: 15px;
57     text-align: center;
58     transition: all 0.3s ease;
59     box-shadow: 0 8px 25px rgba(102, 126, 234, 0.3);
60 }
61
62 .metric-card:hover {
63     transform: scale(1.05);
64     box-shadow: 0 12px 35px rgba(102, 126, 234, 0.4);
65 }
66
67 /* Tab Styles */
68
69 /* Tab Headers */
70
71 .stabs [data-baseweb="tab-list"] {
72     gap: 8px;
73     background: transparent;
74 }
75
76 .stabs [data-baseweb="tab"] {
77     background: #fff;
78     padding: 10px 10px 0 0;
79     padding: 10px 20px;
80     border: none;
81     font-weight: 600;
82     transition: all 0.3s ease;
83     box-shadow: 0 2px 10px rgba(0,0,0,0.1);
84 }
85
86 .stabs [data-baseweb="tab"]:hover {
87     background: #f667eaa;
88     color: #fff;
89     transform: translateY(-2px);
90 }
91
92 .stabs [aria-selected="true"] {
93     background: linear-gradient(135deg, #f667eaa 0%, #f764ba2 100%) !important;
94     color: #fff !important;
95 }
96
97 /* Button Styles */
98 .button {
99     background: linear-gradient(135deg, #f667eaa 0%, #f764ba2 100%) !important;
100     border: none;
101     padding: 10px 24px;
102     border-radius: 10px;
103     font-weight: 600;
104     transition: all 0.3s ease;
105     box-shadow: 0 5px 15px rgba(102, 126, 234, 0.3);
106 }
107
108 .button button:hover {
109     transform: translateY(-2px);
110     box-shadow: 0 8px 25px rgba(102, 126, 234, 0.4);
111     background: linear-gradient(135deg, #f667eaa 0%, #f764ba2 100%) !important;
112 }
113
114 .download-btn {
115     background: linear-gradient(135deg, #f667eaa 0%, #f764ba2 100%) !important;
116 }
117
118 .download-btn:hover {
119     background: linear-gradient(135deg, #f667eaa 0%, #f764ba2 100%) !important;
120 }
121
122 /* Upload Area */
123 .uploaded-area {
124     border: 2px dashed #f667eaa;
125     border-radius: 15px;
126     padding: 3rem;
127     text-align: center;
128     color: #fff;
129     border: 1px solid #f667eaa;
130     transition: all 0.3s ease;
131     margin: 2rem 0;
132 }
```

Figure 36: CSS for Enhance Layout

Activity 4.2: Multi-Tab Application Design

- Tabs implemented:
 1. Dashboard
 2. AI Evaluation
 3. Skills Analysis
 4. Transcript Viewer

```

56
57     # Display results if analysis is complete
58     if st.session_state.analysis_complete and st.session_state.analysis_results:
59         results = st.session_state.analysis_results
60
61     tab1, tab2, tab3, tab4 = st.tabs([
62         "Dashboard", "AI Evaluation", "Skills Analysis", "Transcript"
63     ])

```

Figure 37: Multi-Tab Option

- Each tab displays structured components and analytics.

```

193     with tab3:
194         st.markdown("""
195             <div class="custom-card" style="width: 100%; text-align: center;">
196                 <h3 style="color: #667eea; text-align: center; margin: 0;">
197                     Skills & Qualifications Analysis
198                 </h3>
199             </div>
200             """, unsafe_allow_html=True)
201
202         skills_info = results['skills_info']
203         col1, col2 = st.columns(2)
204
205         with col1:
206             st.subheader("Skill Summary:")
207             if skills_info['skills']:
208                 skills_html = ' '.join([f'<span class="skill-item">{skill}</span>' for skill in skills_info['skills']])
209                 st.markdown(skills_html, unsafe_allow_html=True)
210             else:
211                 st.info("No technical skills detected")
212                 st.markdown("</div>", unsafe_allow_html=True)
213
214         with col2:
215             st.subheader("Languages:")
216             if skills_info['languages']:
217                 languages_html = ' '.join([f'<span class="skill-item">{lang}</span>' for lang in skills_info['languages']])
218                 st.markdown(languages_html, unsafe_allow_html=True)
219             else:
220                 st.info("No languages detected")
221                 st.markdown("</div>", unsafe_allow_html=True)
222
223         col1, col2 = st.columns(2)
224         with col1:
225             st.subheader("Tools & Software:")
226             if skills_info['tools']:
227                 tools_html = ' '.join([f'<span class="skill-item">{tool}</span>' for tool in skills_info['tools']])
228                 st.markdown(tools_html, unsafe_allow_html=True)
229             else:
230                 st.info("No tools detected")
231                 st.markdown("</div>", unsafe_allow_html=True)
232
233         with col2:
234             st.subheader("Education:")
235             if skills_info['degrees']:
236                 degrees_html = ' '.join([f'<span class="skill-item">{degree}</span>' for degree in skills_info['degrees']])
237                 st.markdown(degrees_html, unsafe_allow_html=True)
238             else:
239                 st.info("No education information detected")
240                 st.markdown("</div>", unsafe_allow_html=True)
241                 st.markdown("</div>", unsafe_allow_html=True)

```

Figure 38: Tab - Skill & Qualification Analysis

```

243     with tab2:
244         st.markdown('<div class="fade-in">', unsafe_allow_html=True)
245         st.markdown("""
246             <div class="custom-card" style="width: 100%; text-align: center;">
247                 <h3 style="color: #667eea; text-align: center; margin: 0;">
248                     AI Evaluation & Recommendation
249                 </h3>
250             </div>
251         """), unsafe_allow_html=True)
252
253         st.text_area(
254             "Evaluation Details",
255             results['evaluation'],
256             height=600,
257             key="evaluation_display",
258             label_visibility="collapsed"
259         )
260         st.markdown("</div>", unsafe_allow_html=True)
261         st.markdown("</div>", unsafe_allow_html=True)
262

```

Figure 39: Tab - AI Evaluation & Recommendation

```

168     with tab4:
169         st.markdown("""
170             <div class="custom-card" style="width: 100%; text-align: center;">
171                 <h3 style="color: #667eea; text-align: center; margin: 0;">
172                     Conversation Transcript
173                 </h3>
174             </div>
175         """), unsafe_allow_html=True)
176         st.markdown(**Speaker Legend:** ● Candidate | ● Interviewer")
177
178         for segment in results['diarized_segments']:
179             speaker_class = "Candidate" if segment['speaker'] == results['candidate_speaker'] else "Interviewer"
180             speaker_emoji = "●" if segment['speaker'] == results['candidate_speaker'] else "●"
181
182             st.markdown(f"""
183                 <div class="transcript-segment {speaker_class}">
184                     <div style="display: flex; justify-content: space-between; align-items: center; margin-bottom: 0.5rem;">
185                         <strong>{speaker_emoji} {speaker_class}</strong> : &nbsp;
186                         <small style="color: #666;">{segment['start']:.1f}s - {segment['end']:.1f}s</small>
187                     </div>
188                     <div style="color: #333;">{segment['text']}</div>
189                 """), unsafe_allow_html=True)
190
191         st.markdown("</div>", unsafe_allow_html=True)
192

```

Figure 40: Tab - Conversation Transcript

```

115     with tab1:
116         st.markdown('<div class="fade-in">', unsafe_allow_html=True)
117         st.markdown("""
118             <div class="custom-card" style="width: 100%; text-align: center;">
119                 <h3 style="color: #667eea; text-align: center; margin: 0;">
120                     Interview Overview Dashboard
121                 </h3>
122             </div>
123         "", unsafe_allow_html=True)
124
125         col1, col2, col3, col4 = st.columns(4)
126
127         with col1:
128             st.markdown(f"""
129                 <div class="metric-card">
130                     <h3>Duration</h3>
131                     <h2>{results['audio_metadata'][0]['duration']:.1f}</h2>
132                     <small>seconds</small>
133                 </div>
134             "", unsafe_allow_html=True)
135
136         with col2:
137             sentiment = results['sentiment']
138             st.markdown(f"""
139                 <div class="metric-card">
140                     <h3>Sentiment</h3>
141                     <h2>{sentiment['label'].title()}</h2>
142                     <small>Score: {sentiment['score']:.2f}</small>
143                 </div>
144             "", unsafe_allow_html=True)
145
146         with col3:
147             skills_count = len(results['skills_info']['skills'])
148             st.markdown(f"""
149                 <div class="metric-card">
150                     <h3>Skills</h3>
151                     <h2>{skills_count}</h2>
152                     <small>Detected</small>
153                 </div>
154             "", unsafe_allow_html=True)
155
156         with col4:
157             segments_count = len(results['diarized_segments'])
158             st.markdown(f"""
159                 <div class="metric-card">
160                     <h3>Segments</h3>
161                     <h2>{segments_count}</h2>
162                     <small>Conversation</small>
163                 </div>
164             "", unsafe_allow_html=True)
165
166         st.markdown("</div>", unsafe_allow_html=True)
167

```

Figure 41: Tab - Interview Overview Dashboard

Activity 4.3: Audio Upload & Processing Workflow

- Implement drag-and-drop audio uploader.

```

59     # File upload section
60     st.markdown("""
61         <div class="upload-area fade-in">
62             <h3 style="color: #667eea; margin-bottom: 1rem;">Upload Interview Audio</h3>
63             <p style="color: #666; margin-bottom: 2rem;">Supported formats: MP3, WAV, M4A</p>
64         </div>
65     "", unsafe_allow_html=True)
66
67     uploaded_file = st.file_uploader(
68         "Choose an audio file",
69         type=['mp3', 'wav', 'm4a'],
70         help="Upload your interview recording for analysis",
71         label_visibility="collapsed"
72     )
73
74     # Store uploaded file in session state
75     if uploaded_file is not None:
76         st.session_state.uploaded_file = uploaded_file
77         st.success(f"File ready: {uploaded_file.name}")
78
79     st.markdown(" ")
80

```

Figure 42: Uploading Audio

- Provide progress indicators (0–100%) during analysis with display success or error messages.

```
81     if st.session_state.uploaded_file is not None and not st.session_state.analysis_complete:
82         if st.button("Start AI Analysis", type="primary", use_container_width=True):
83
84             with tempfile.NamedTemporaryFile(delete=False, suffix=os.path.splitext(st.session_state.uploaded_file.name)[1]) as tmp_file:
85                 tmp_file.write(st.session_state.uploaded_file.getvalue())
86                 audio_path = tmp_file.name
87
88         try:
89             with st.spinner("AI is analyzing your interview... This may take 5-7 minutes."):
90                 progress_bar = st.progress(0)
91
92                 progress_bar.progress(50)
93                 results = run_full_pipeline(audio_path)
94                 st.session_state.analysis_results = results
95                 st.session_state.analysis_complete = True
96                 progress_bar.progress(100)
97
98             st.success("Analysis completed successfully!")
99             st.rerun()
100
101         except Exception as e:
102             st.error(f"Error analyzing interview: {str(e)}")
103         finally:
104             if os.path.exists(audio_path):
105                 os.unlink(audio_path)
```

Figure 43: Progress Bar

Activity 4.4: Skills, Sentiment & Transcript Visualization

- Skill and other metrics display with custom styling.

```
115
116     with tab1:
117         st.markdown('<div class="fade-in">', unsafe_allow_html=True)
118         st.markdown("""
119             <div class="custom-card" style="width: 100%; text-align: center;">
120                 <h3 style="color: #6677ee; text-align: center; margin: 0;">
121                     Interview Overview Dashboard
122                 </h3>
123             </div>
124             """, unsafe_allow_html=True)
125
126         col1, col2, col3, col4 = st.columns(4)
127
128     with col1:
129         st.markdown(f"""
130             <div class="metric-card">
131                 <h3>Duration</h3>
132                 <h2>{results['audio_metadata']['duration']:.1f}</h2>
133                 <small>seconds</small>
134             </div>
135             """, unsafe_allow_html=True)
136
137     with col2:
138         sentiment = results['sentiment']
139         st.markdown(f"""
140             <div class="metric-card">
141                 <h3>Sentiment</h3>
142                 <h2>{sentiment['label'].title()}</h2>
143                 <small>Score: {sentiment['score']:.2f}</small>
144             </div>
145             """, unsafe_allow_html=True)
146
147     with col3:
148         skills_count = len(results['skills_info']['skills'])
149         st.markdown(f"""
150             <div class="metric-card">
151                 <h3>Skills</h3>
152                 <h2>{skills_count}</h2>
153                 <small>Detected</small>
154             </div>
155             """, unsafe_allow_html=True)
156
157     with col4:
158         segments_count = len(results['diarized_segments'])
159         st.markdown(f"""
160             <div class="metric-card">
161                 <h3>Segments</h3>
162                 <h2>{segments_count}</h2>
163                 <small>Conversation</small>
164             </div>
165             """, unsafe_allow_html=True)
166
167         st.markdown("</div>", unsafe_allow_html=True)
```

Figure 44: Visualization of tabs

- Transcript with speaker colour coding.

```

168     with tab4:
169         st.markdown("""
170             <div class="custom-card" style="width: 100%; text-align: center;">
171                 <h3 style="color: #667eea; text-align: center; margin: 0;">
172                     | Conversation Transcript
173                 </h3>
174             </div>
175             """", unsafe_allow_html=True)
176         st.markdown("***Speaker Legend:** 🟦 Candidate | 🟧 Interviewer")
177
178         for segment in results['diarized_segments']:
179             speaker_class = "Candidate" if segment['speaker'] == results['candidate_speaker'] else "Interviewer"
180             speaker_emoji = "🟡" if segment['speaker'] == results['candidate_speaker'] else "🔴"
181
182             st.markdown(f"""
183                 <div class="transcript-segment {speaker_class}">
184                     <div style="display: flex; justify-content: space-between; align-items: center; margin-bottom: 0.5rem;">
185                         <strong>{speaker_emoji} {speaker_class}</strong>
186                         <small style="color: #666; margin-left: 10px;">{segment['start']:.1f}s - {segment['end']:.1f}s</small>
187                     </div>
188                     <div style="color: #333; margin-top: 10px; font-family: monospace; padding: 5px; border-radius: 5px; background-color: #f9f9f9; border: 1px solid #ccc; width: fit-content; margin-left: auto;">{segment['text']}

```

Figure 45: Custom colour for Speakers

Activity 4.5: Report Export System

- Generate downloadable reports in:
 - TXT
 - JSON
 - PDF

```

263     # Download Section
264     st.markdown("----")
265     st.markdown("""
266         <div class="fade-in">
267             <h2>Download Professional Reports</h2>
268             <p>Export comprehensive analysis in your preferred format</p>
269         </div>
270     """", unsafe_allow_html=True)
271
272     timestamp = datetime.now().strftime("%Y-%m-%d %H:%M:%S")
273     default_name = f"professional_interview_analysis_{timestamp}"
274
275     col1, col2, col3 = st.columns(3)
276
277     with col1:
278         txt_data = ""
279         try:
280             with tempfile.NamedTemporaryFile(mode='w', delete=False, suffix='.txt') as tmp_file:
281                 export_txt(results, tmp_file.name)
282                 with open(tmp_file.name, 'r', encoding='utf-8') as f:
283                     txt_data = f.read()
284                 os.unlink(tmp_file.name)
285         except Exception as e:
286             txt_data = f"Error generating TXT: {str(e)}"
287
288         st.download_button(
289             label="Download TXT Report",
290             data=txt_data,
291             file_name=f"{default_name}.txt",
292             mime="text/plain",
293             use_container_width=True,
294             key="download_txt"
295         )
296
297     with col2:
298         json_data = ""
299         try:
300             with tempfile.NamedTemporaryFile(mode='w', delete=False, suffix='.json') as tmp_file:
301                 export_json(results, tmp_file.name)
302                 with open(tmp_file.name, 'r', encoding='utf-8') as f:
303                     json_data = f.read()
304                 os.unlink(tmp_file.name)
305         except Exception as e:
306             json_data = json.dumps({"error": f"Error generating JSON: {str(e)}"})
307
308         st.download_button(
309             label="Download JSON Report",
310             data=json_data,
311             file_name=f"{default_name}.json",
312             mime="application/json",
313             use_container_width=True,
314             key="download_json"
315         )

```

Figure 46: Download button for each format

```
317     with col3:
318         pdf_data = b""
319         try:
320             with tempfile.NamedTemporaryFile(delete=False, suffix=".pdf") as tmp_file:
321                 export_pdf(results, tmp_file.name)
322                 with open(tmp_file.name, 'rb') as f:
323                     pdf_data = f.read()
324                 os.unlink(tmp_file.name)
325         except Exception as e:
326             try:
327                 with tempfile.NamedTemporaryFile(mode='w', delete=False, suffix=".txt") as tmp_file:
328                     export_txt(results, tmp_file.name)
329                     with open(tmp_file.name, 'r', encoding='utf-8') as f:
330                         pdf_data = f.read().encode('utf-8')
331                     os.unlink(tmp_file.name)
332             except:
333                 pdf_data = f"Error generating PDF: {str(e)}".encode('utf-8')
334
335         st.download_button(
336             label="Download PDF Report",
337             data=pdf_data,
338             file_name=f'{default_name}.pdf',
339             mime="application/pdf",
340             use_container_width=True,
341             key="download_pdf"
342         )
```

Figure 47: Download Button

MILESTONE 5: Testing, Optimization & Deployment

This milestone ensures that the AI Interview Transcript Analyzer delivers a polished, intuitive, and fully functional user experience. The focus is on visually validating UI components, testing the display of interview insights, and confirming that users can seamlessly upload audio, view analysis results, and download reports. Performance checks and deployment finalization are handled afterward.

Activity 5.1: UI Display Testing & Result Verification

- **Validate Home Page UI and Feature Cards**

Verify that feature sections (Transcription, Diarization, Sentiment, Skills) render with proper styling.



Figure 48: Streamlit Home Page

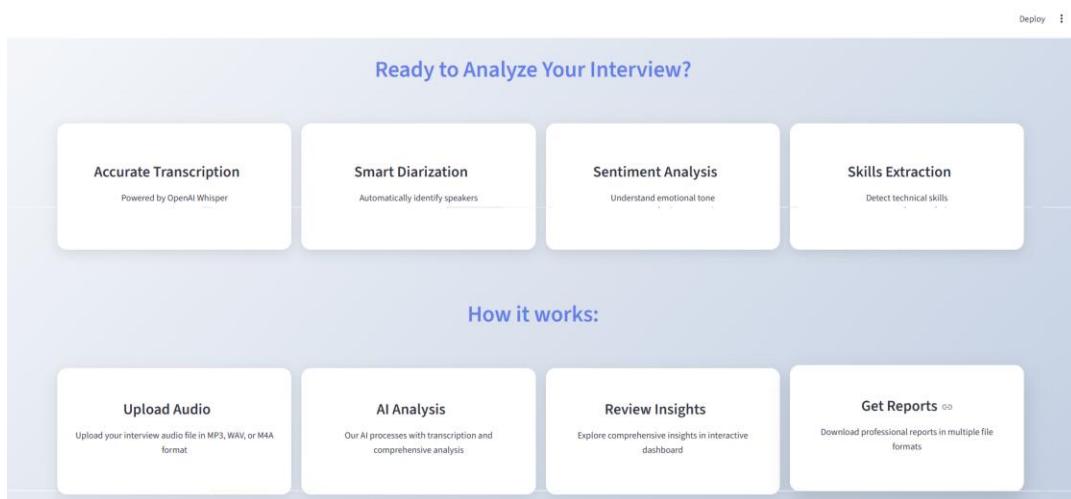


Figure 49: Streamlit About Section

- **Test Audio Upload Workflow**

Upload MP3/WAV/M4A files and confirm success messages appear.

Verify file handler stores uploaded audio in session state.

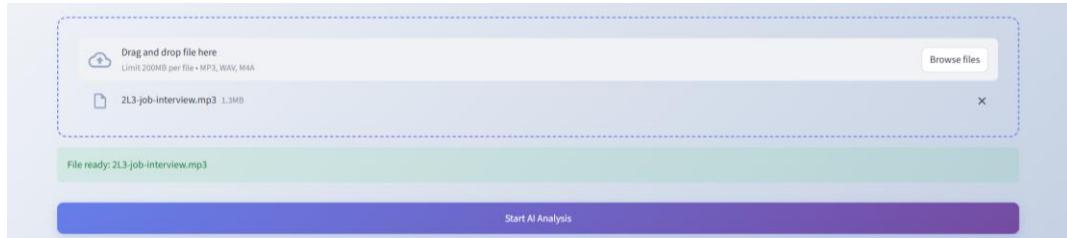


Figure 50: Uploading Section

- **Verify Processing Feedback & Progress Bar**

Validate progress indicators during analysis (0% → 50% → 100%).

Confirm loading animations and spinner messages appear.

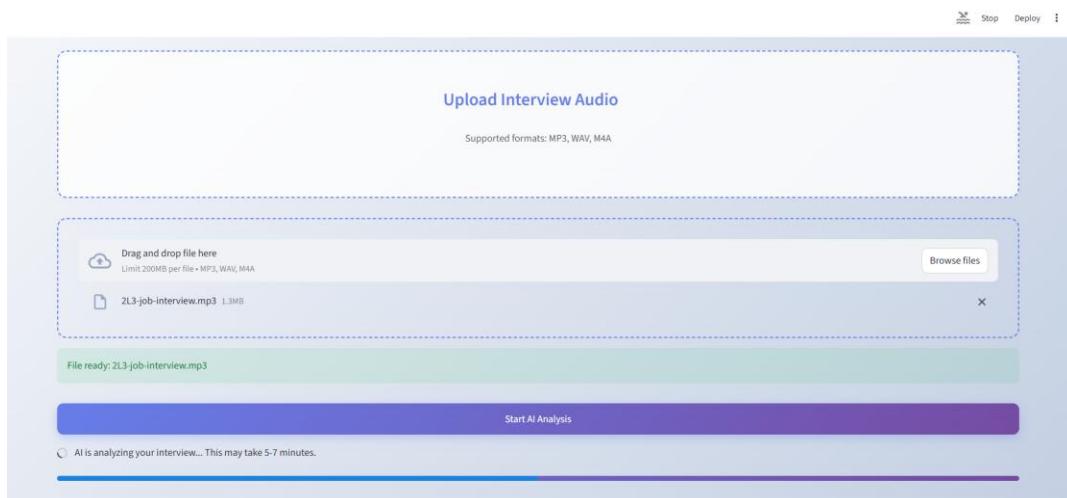


Figure 51: Progress Bar

- **Dashboard Metrics Display Testing**

Check that Duration, Sentiment Score, Skill Count, and Segments display correctly in metric cards.

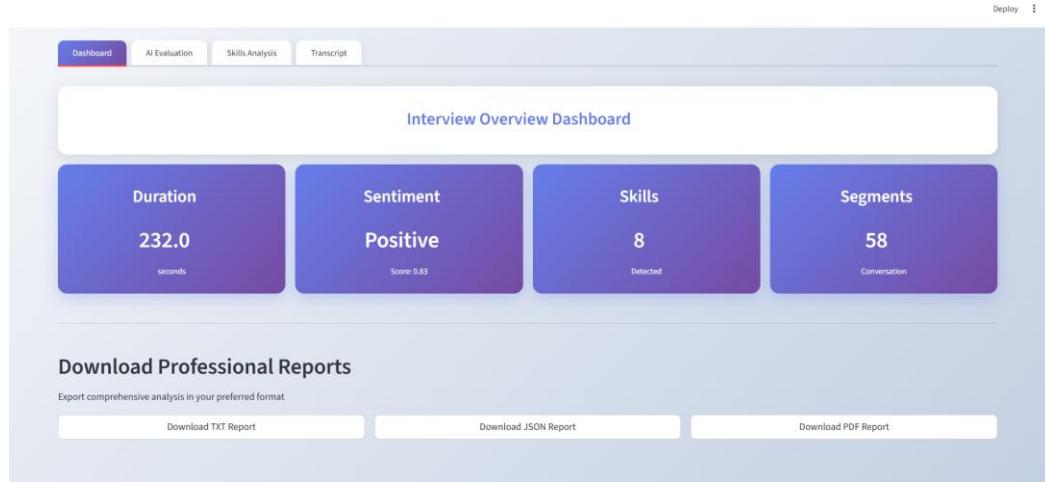


Figure 52: Dashboard

- **Transcript Viewer Validation**

Confirm diarized segments render with Candidate/Interviewer colour coding.

Validate timestamp formatting and segment spacing.

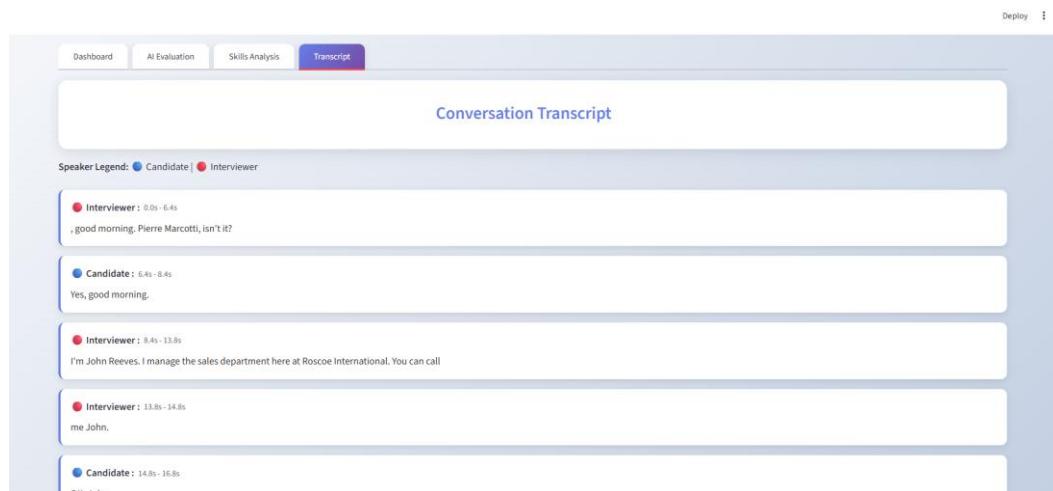


Figure 53: Transcript

- **Skills Analysis Tab UI Testing**

Check tag-based display for Skills, Languages, Tools, and Education.

Validate responsive grid layout for different screen sizes.

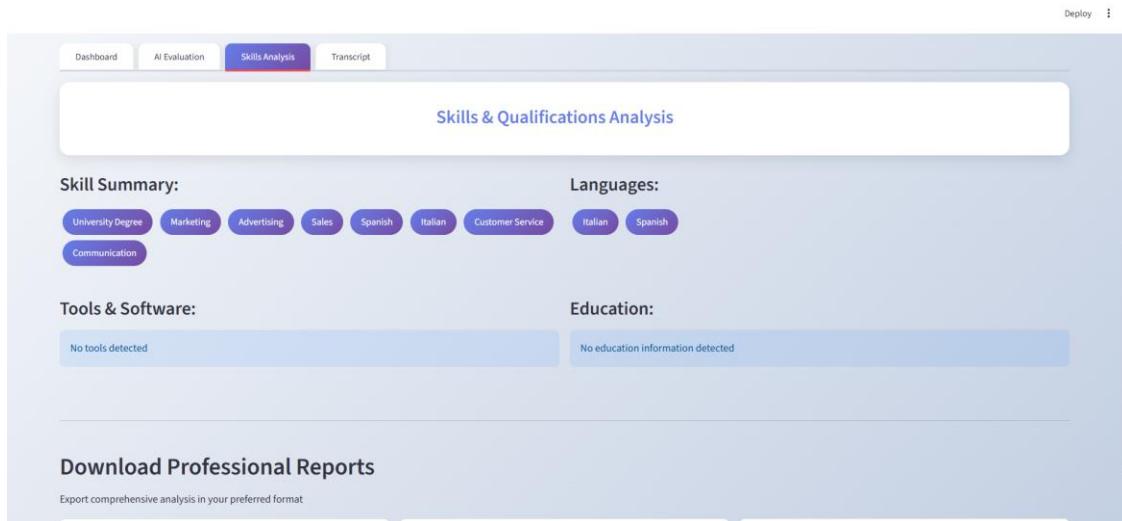


Figure 54: Skill Extractor

- **AI Evaluation Display Testing**

Verify evaluation text loads fully inside scrollable text area.

Check formatting of Summary, Recommendation, Sections, and Reasoning.

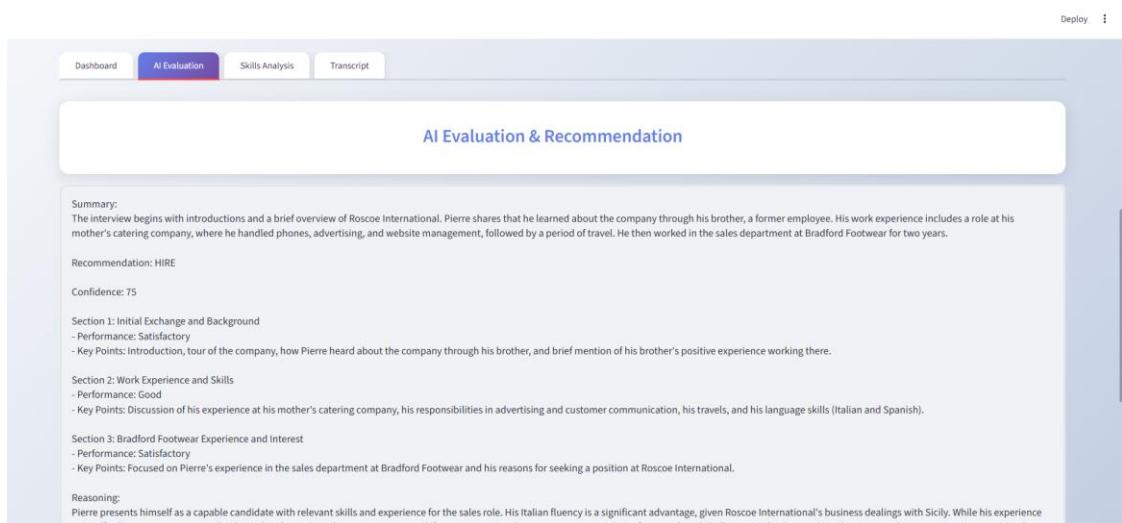


Figure 55: AI Evaluation

- **Download Buttons Testing**

Ensure TXT, JSON, PDF buttons appear correctly with icons/styling.

Verify downloaded files have correct names and valid content.

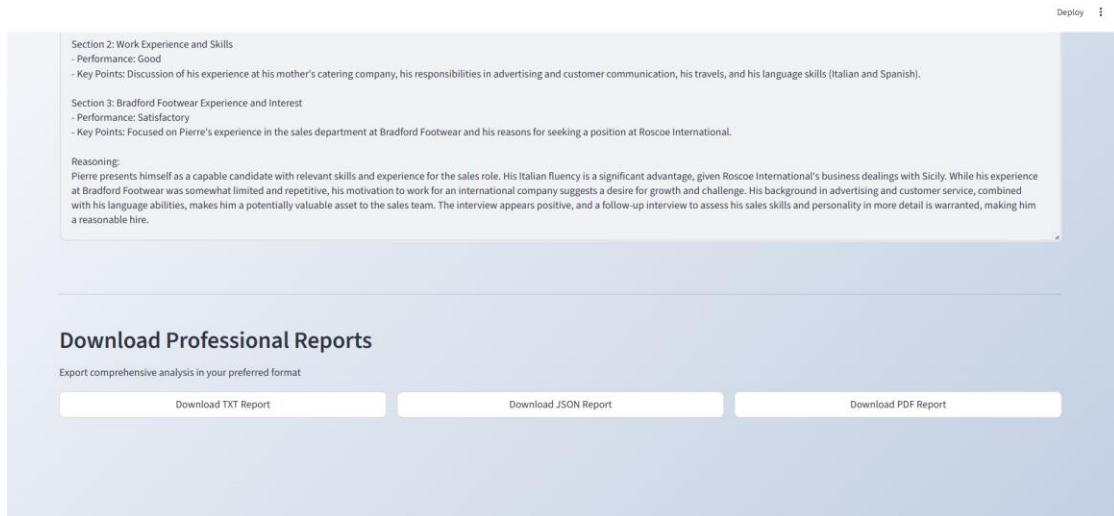


Figure 56: Download Buttons

Activity 5.2: Deployment Preparation & Validation

- **Prepare Streamlit for Production Deployment**

Configured Streamlit page settings, optimized custom CSS loading, and validated UI components to ensure smooth rendering in the deployed environment.

- **Verify Environment Variables and Model Access**

Confirmed the GEMINI_API_KEY and all model configurations load correctly on the server, ensuring Whisper, Resemblyzer, and HuggingFace models function without errors.

- **Run End-to-End Tests on Deployed Application**

Executed the full analysis pipeline online upload, transcription, diarization, sentiment, skills, evaluation, and report export to ensure stable performance across browsers and devices.

Conclusion

The AI Interview Transcript Analyzer represents a comprehensive, intelligent solution that transforms traditional interview evaluation into an automated, data-driven, and highly efficient process. By integrating advanced audio processing, natural language understanding, sentiment modelling, and Gemini-powered reasoning, the system delivers deep insights that would otherwise require extensive manual effort from HR teams. From transcription and speaker diarization to skill extraction and final hiring recommendations, each component works in harmony to provide a complete overview of candidate performance.

Through its modular architecture and interactive Streamlit interface, the platform ensures both technical robustness and exceptional user experience. Recruiters, trainers, and organizations can seamlessly upload audio, explore visually rich insights across dashboards, evaluate candidate strengths, and export professional reports all within a streamlined workflow. The design emphasizes accuracy, transparency, and usability, making it suitable for real recruitment, interview practice, and organizational assessment systems.

Looking forward, the platform provides a strong foundation for enhancements such as multilingual transcription, advanced communication analytics, behavioral scoring, ATS integration, and voice sentiment patterns. With its scalable architecture and expandable AI capabilities, the AI Interview Transcript Analyzer demonstrates how modern AI technologies can revolutionize hiring processes, improve decision-making quality, and set new standards for interview intelligence.